

Convolutional Neural Networks

Presented By: Don Kim - 2022.12.11

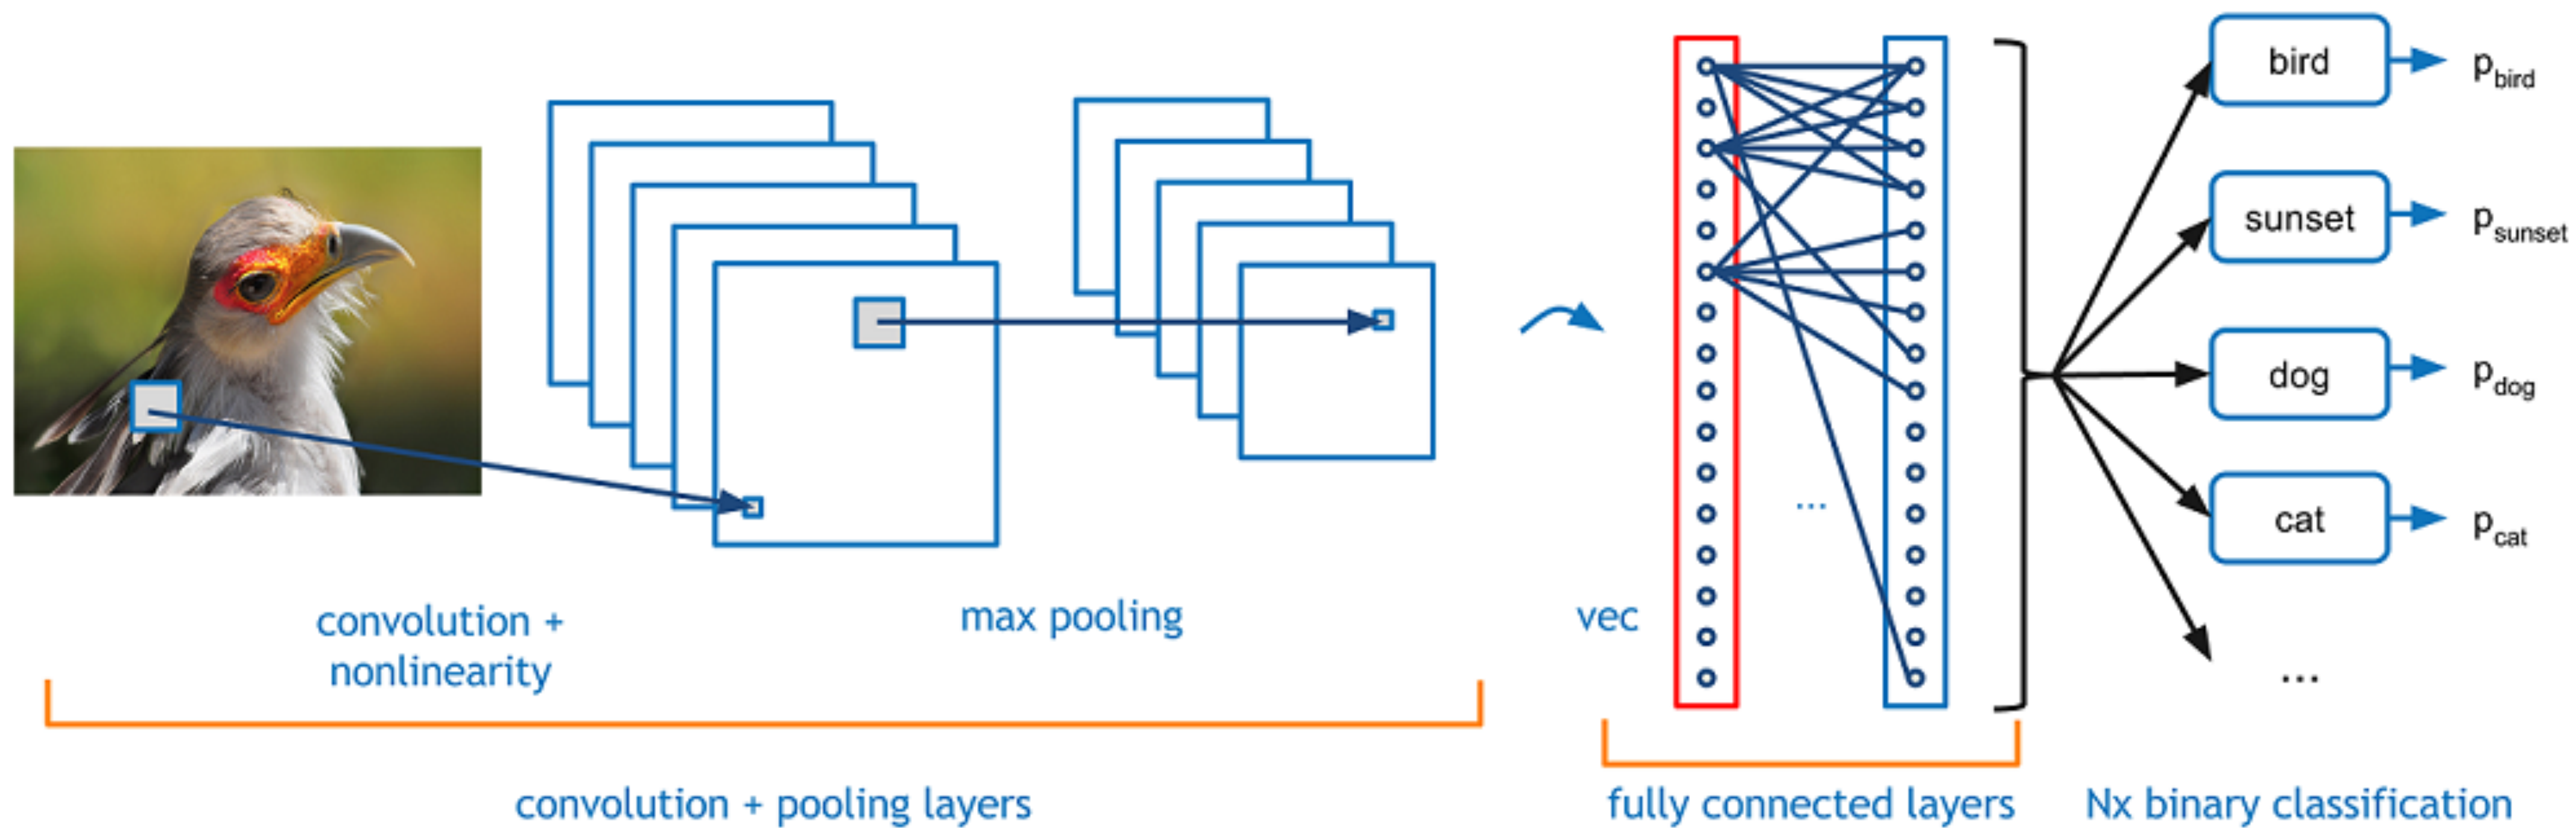
Project Description

- 1. Implement CNN in the GWU Neural Network Library**
- 2. Determine the effectiveness of the implementation with a test case**
- 3. Cannot use any existing CNN libraries (E.g. Tensorflow)**

Convolutional Neural Network (CNN)

What is a CNN?

Deep learning algorithm that takes in an **image** as an input, assigns importance (weight) to certain features in the image, and is able to **classify the image input**



CNNs are great for image classification compared to other models, since the **convolution layer** reduces images without losing its information

Parts of the CNN

- 1. Input Layer**
- 2. Convolutional Layer**
- 3. Max Pooling Layer**
- 4. Flattening Layer**
- 5. Dense Layer**
- 6. Activation / Loss Layer**
- 7. Output**

For the project, the Convolutional, Max Pooling, and Flattening layer needed to be implemented

```
network = GWUNetwork()  
network.add(Conv2D(...))  
network.add(MaxPooling2D(...))  
network.add(Flatten(...))  
network.add(Dense(...))  
network.add(Dense(...))
```

Convolutional Layer

“Sliding a filter over an image”

Purpose: Extracts the high-level features to a feature map, such as edges, from the input image by applying a **kernel** / filter

Kernel = Matrix of size $N \times N$ that moves over the input image and performs dot product with the input image sub-regions (weights set randomly)

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4



*

1	0	-1
1	0	-1
1	0	-1

=

6		

$$\begin{aligned}
 &7 \times 1 + 4 \times 1 + 3 \times 1 + \\
 &2 \times 0 + 5 \times 0 + 3 \times 0 + \\
 &3 \times -1 + 3 \times -1 + 2 \times -1 \\
 &= 6
 \end{aligned}$$



After convolution, of an image size of $M \times M$ and a kernel of size $N \times N$, the resulting dimension is $(M - N) + 1$

E.g. Convolution of 28x28 Image, 3x3 kernel, and “y” filters results in $(28 - 3) + 1 \Rightarrow (y, 24, 24)$ feature map

Forward Propagation

```
output = np.zeros((self.num_filters, convRow, convColumn))

for i in range (self.num_filters):
    for x in range(convRow):
        for y in range (convColumn):
            for z in range(self.kernel_size):
                for v in range (self.kernel_size):
                    output[i, x, y] += input[x + z, y + v] * currentFilters[i, z, v]
```

Input = (1 , 28, 28)

Kernel Size = (3 , 3)

Output = (num_filters, (28 - 3) +1, (28 - 3)+1)

Back Propagation

```
    for x in range (0, self.kernel_size):
        for y in range (0, self.kernel_size):
            for z in range (0, self.convolve_size):
                for v in range (0, self.convolve_size):
                    kernel_gradient[i, x, y] += self.input[x + z, y + v] * output_error[i, z, v]

# Update Filters with Learning Rate
self.kernel -= np.array(kernel_gradient) * learning_rate
```

1. Filter Update

Perform convolution between the **original input into Convolutional layer** and **loss gradient** from previous layer to get “update_values”

Then, multiply “update_values” with the **learning rate** and update the current filter values

$$\begin{array}{|c|c|} \hline \frac{\partial L}{\partial F_{11}} & \frac{\partial L}{\partial F_{12}} \\ \hline \frac{\partial L}{\partial F_{21}} & \frac{\partial L}{\partial F_{22}} \\ \hline \end{array} = \text{Convolution} \left(\begin{array}{|c|c|c|} \hline X_{11} & X_{12} & X_{13} \\ \hline X_{21} & X_{22} & X_{23} \\ \hline X_{31} & X_{32} & X_{33} \\ \hline \end{array}, \begin{array}{|c|c|} \hline \frac{\partial L}{\partial O_{11}} & \frac{\partial L}{\partial O_{12}} \\ \hline \frac{\partial L}{\partial O_{21}} & \frac{\partial L}{\partial O_{22}} \\ \hline \end{array} \right)$$

where

$$\begin{array}{|c|c|c|} \hline X_{11} & X_{12} & X_{13} \\ \hline X_{21} & X_{22} & X_{23} \\ \hline X_{31} & X_{32} & X_{33} \\ \hline \end{array} = \text{Input } X$$

$$\begin{array}{|c|c|} \hline \frac{\partial L}{\partial O_{11}} & \frac{\partial L}{\partial O_{12}} \\ \hline \frac{\partial L}{\partial O_{21}} & \frac{\partial L}{\partial O_{22}} \\ \hline \end{array} = \frac{\partial L}{\partial O} \text{ Loss gradient from previous layer}$$

2. Gradient

Perform **full convolution** between the passed **loss gradient** and **current filter** to get “new_gradient”

Then, pass “new_gradient” into the next layer in back propagation

Backpropagation in a Convolutional Layer of a CNN

Finding the gradients:

$$\frac{\partial L}{\partial F} = \text{Convolution} \left(\text{Input } X, \text{ Loss gradient } \frac{\partial L}{\partial O} \right)$$

$$\frac{\partial L}{\partial X} = \text{Full Convolution} \left(\begin{array}{c} 180^\circ \text{rotated} \\ \text{Filter } F \end{array}, \text{ Loss Gradient } \frac{\partial L}{\partial O} \right)$$

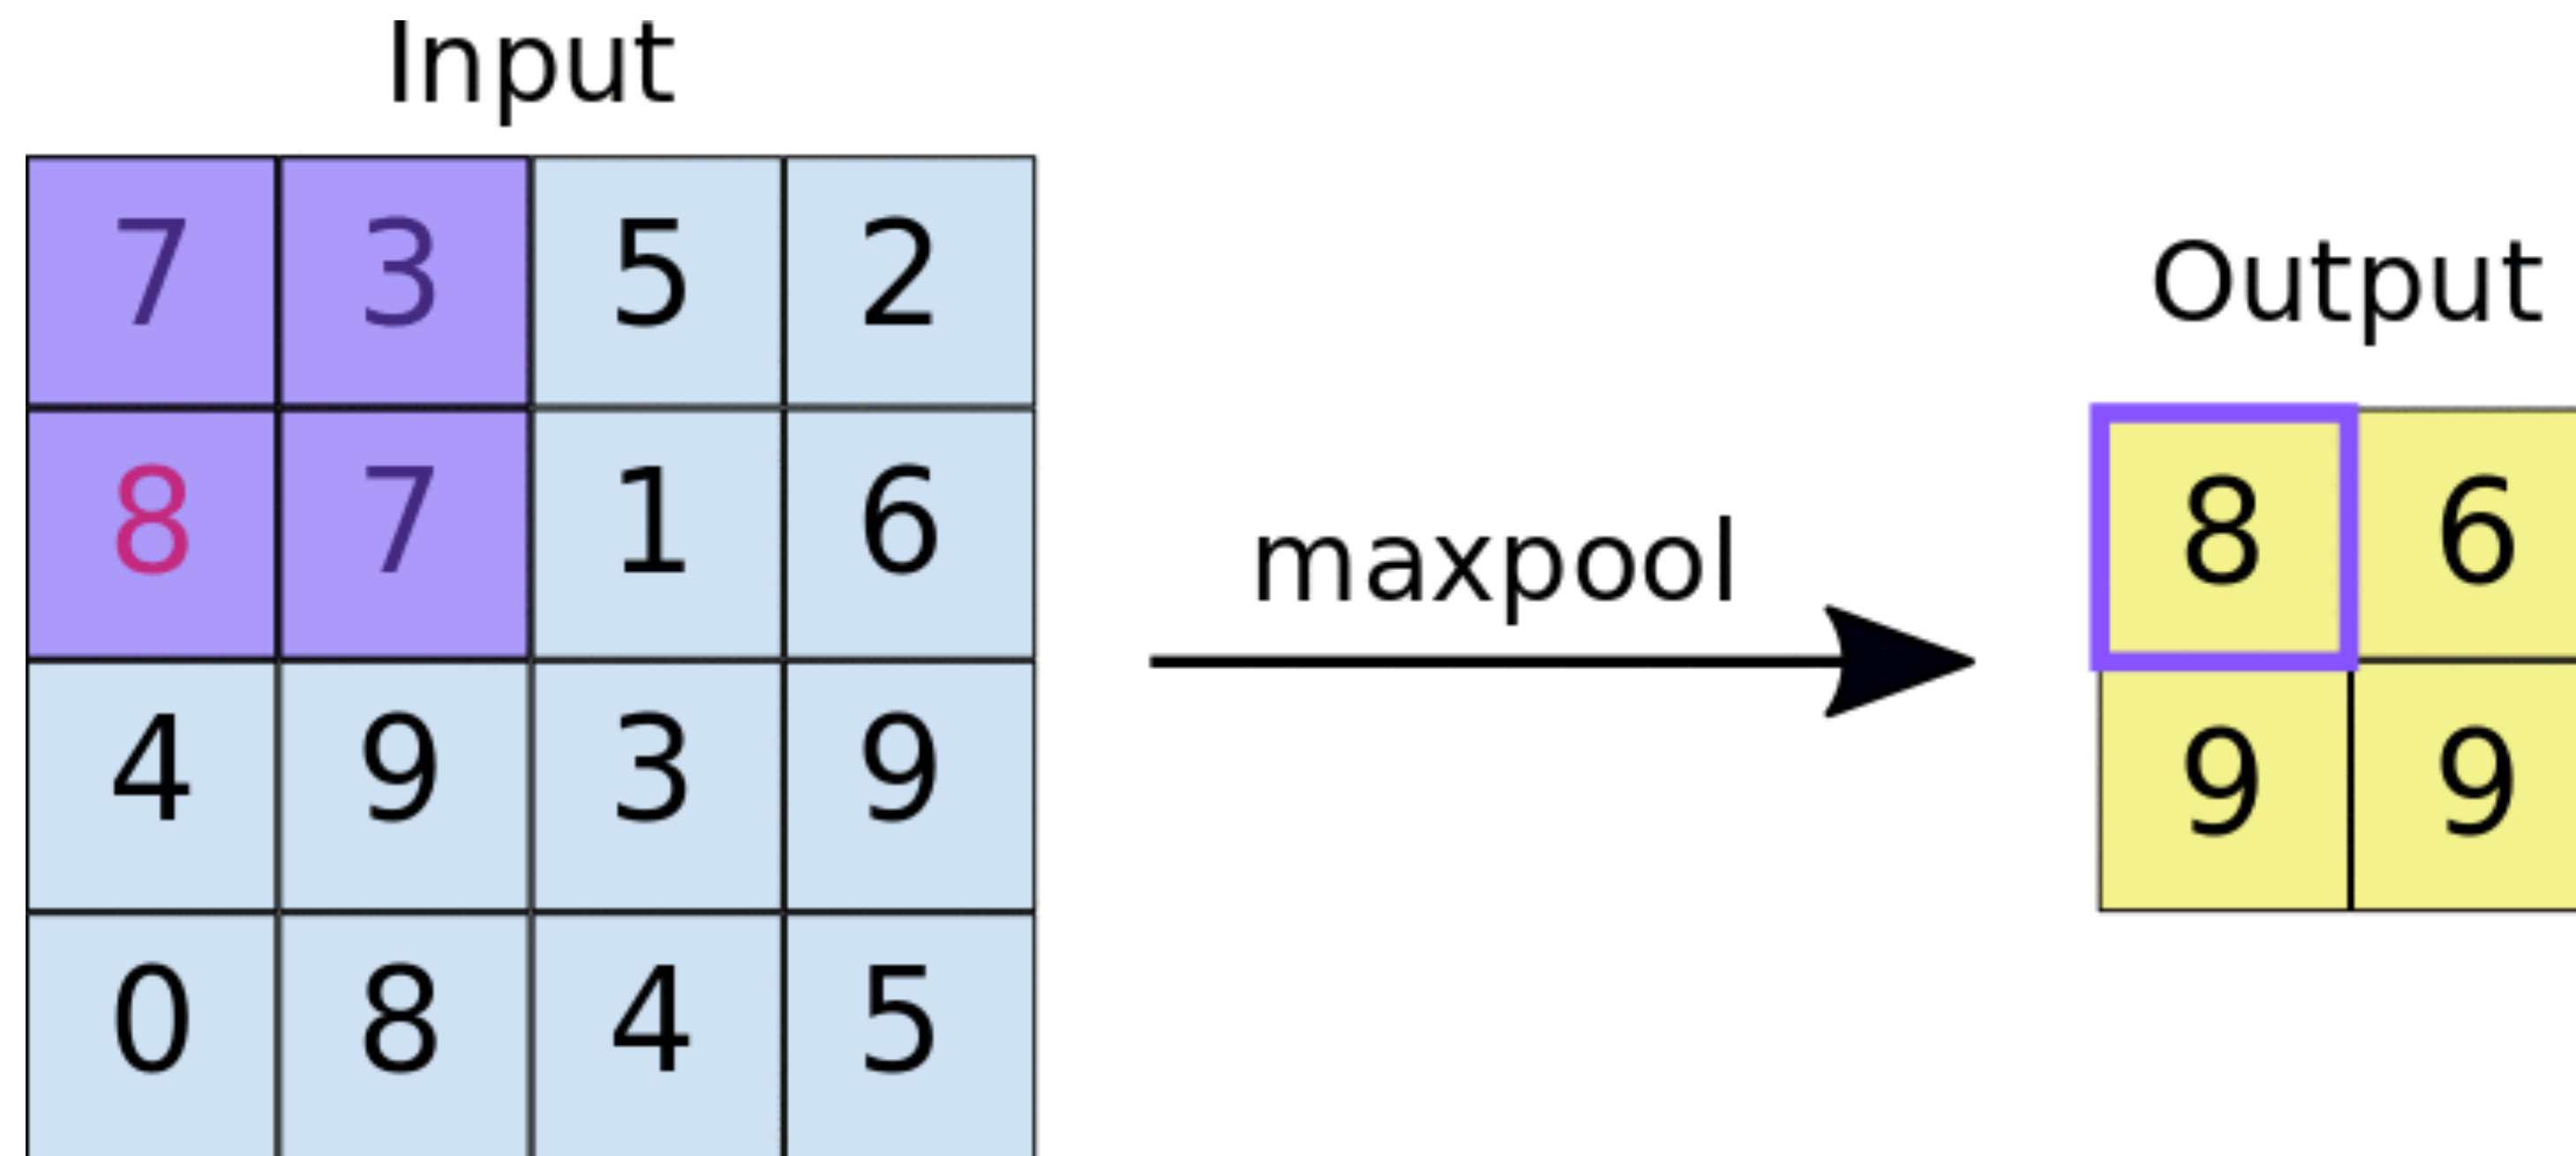
Max Pooling Layer

Purpose: Used to reduce the dimensions of feature maps and reduce the amount of computation required by the network



Moves over the feature maps and selects the **maximum element from the $N \times N$ filter (pooling region)**

Pooling Size = NxN window that moves across the subregions of the feature maps

Strides = Integer that specifies how far the pooling window moves on each step



Pooling Size = 2×2
Strides = 2



Forward Propagation

```
output = np.zeros((num_filters, self.output_size, self.output_size))

for i in range(0, num_filters):
    for x in range(0, tempOutputSize):
        for y in range(0, tempOutputSize):
            tempArray = input[i, x*self.strides:(x*self.strides)+self.pool_size, y*self.strides:(y*self.strides)+self.pool_size]
            output[i, x, y] = np.max(tempArray)
```

$x * \text{self.strides} : (x * \text{self.strides}) + \text{self.pool_size}$

$y * \text{self.strides} : (y * \text{self.strides}) + \text{self.pool_size}$

MaxPooling layer basically **halves** the feature maps

Back Propagation

```
input_gradient = np.zeros(self.input_shape)

for i in range(0, self.num_filters):
    y_coord = 0
    for x in range(0, self.output_size):
        x_coord = 0
        for y in range(0, self.output_size):
            input_sub = self.input[i, x*self.strides:(x*self.strides)+self.pool_size, y*self.strides:(y*self.strides)+self.pool_size]
            max = np.max(input_sub)
            result = unravel_index(input_sub.argmax(), input_sub.shape)
            max_x = result[0]
            max_y = result[1]
            input_gradient[i, x*self.strides:(x*self.strides)+self.pool_size, y*self.strides:(y*self.strides)+self.pool_size][max_x, max_y] = max
        return input_gradient
```

Finds the maximum element from the array within the pooling window and returns an array filled with zeros, except the maximum elements within each window

1	2	4	5
7	9	20	3

1	2
7	9

4	5
20	3

Pooling Size = 2

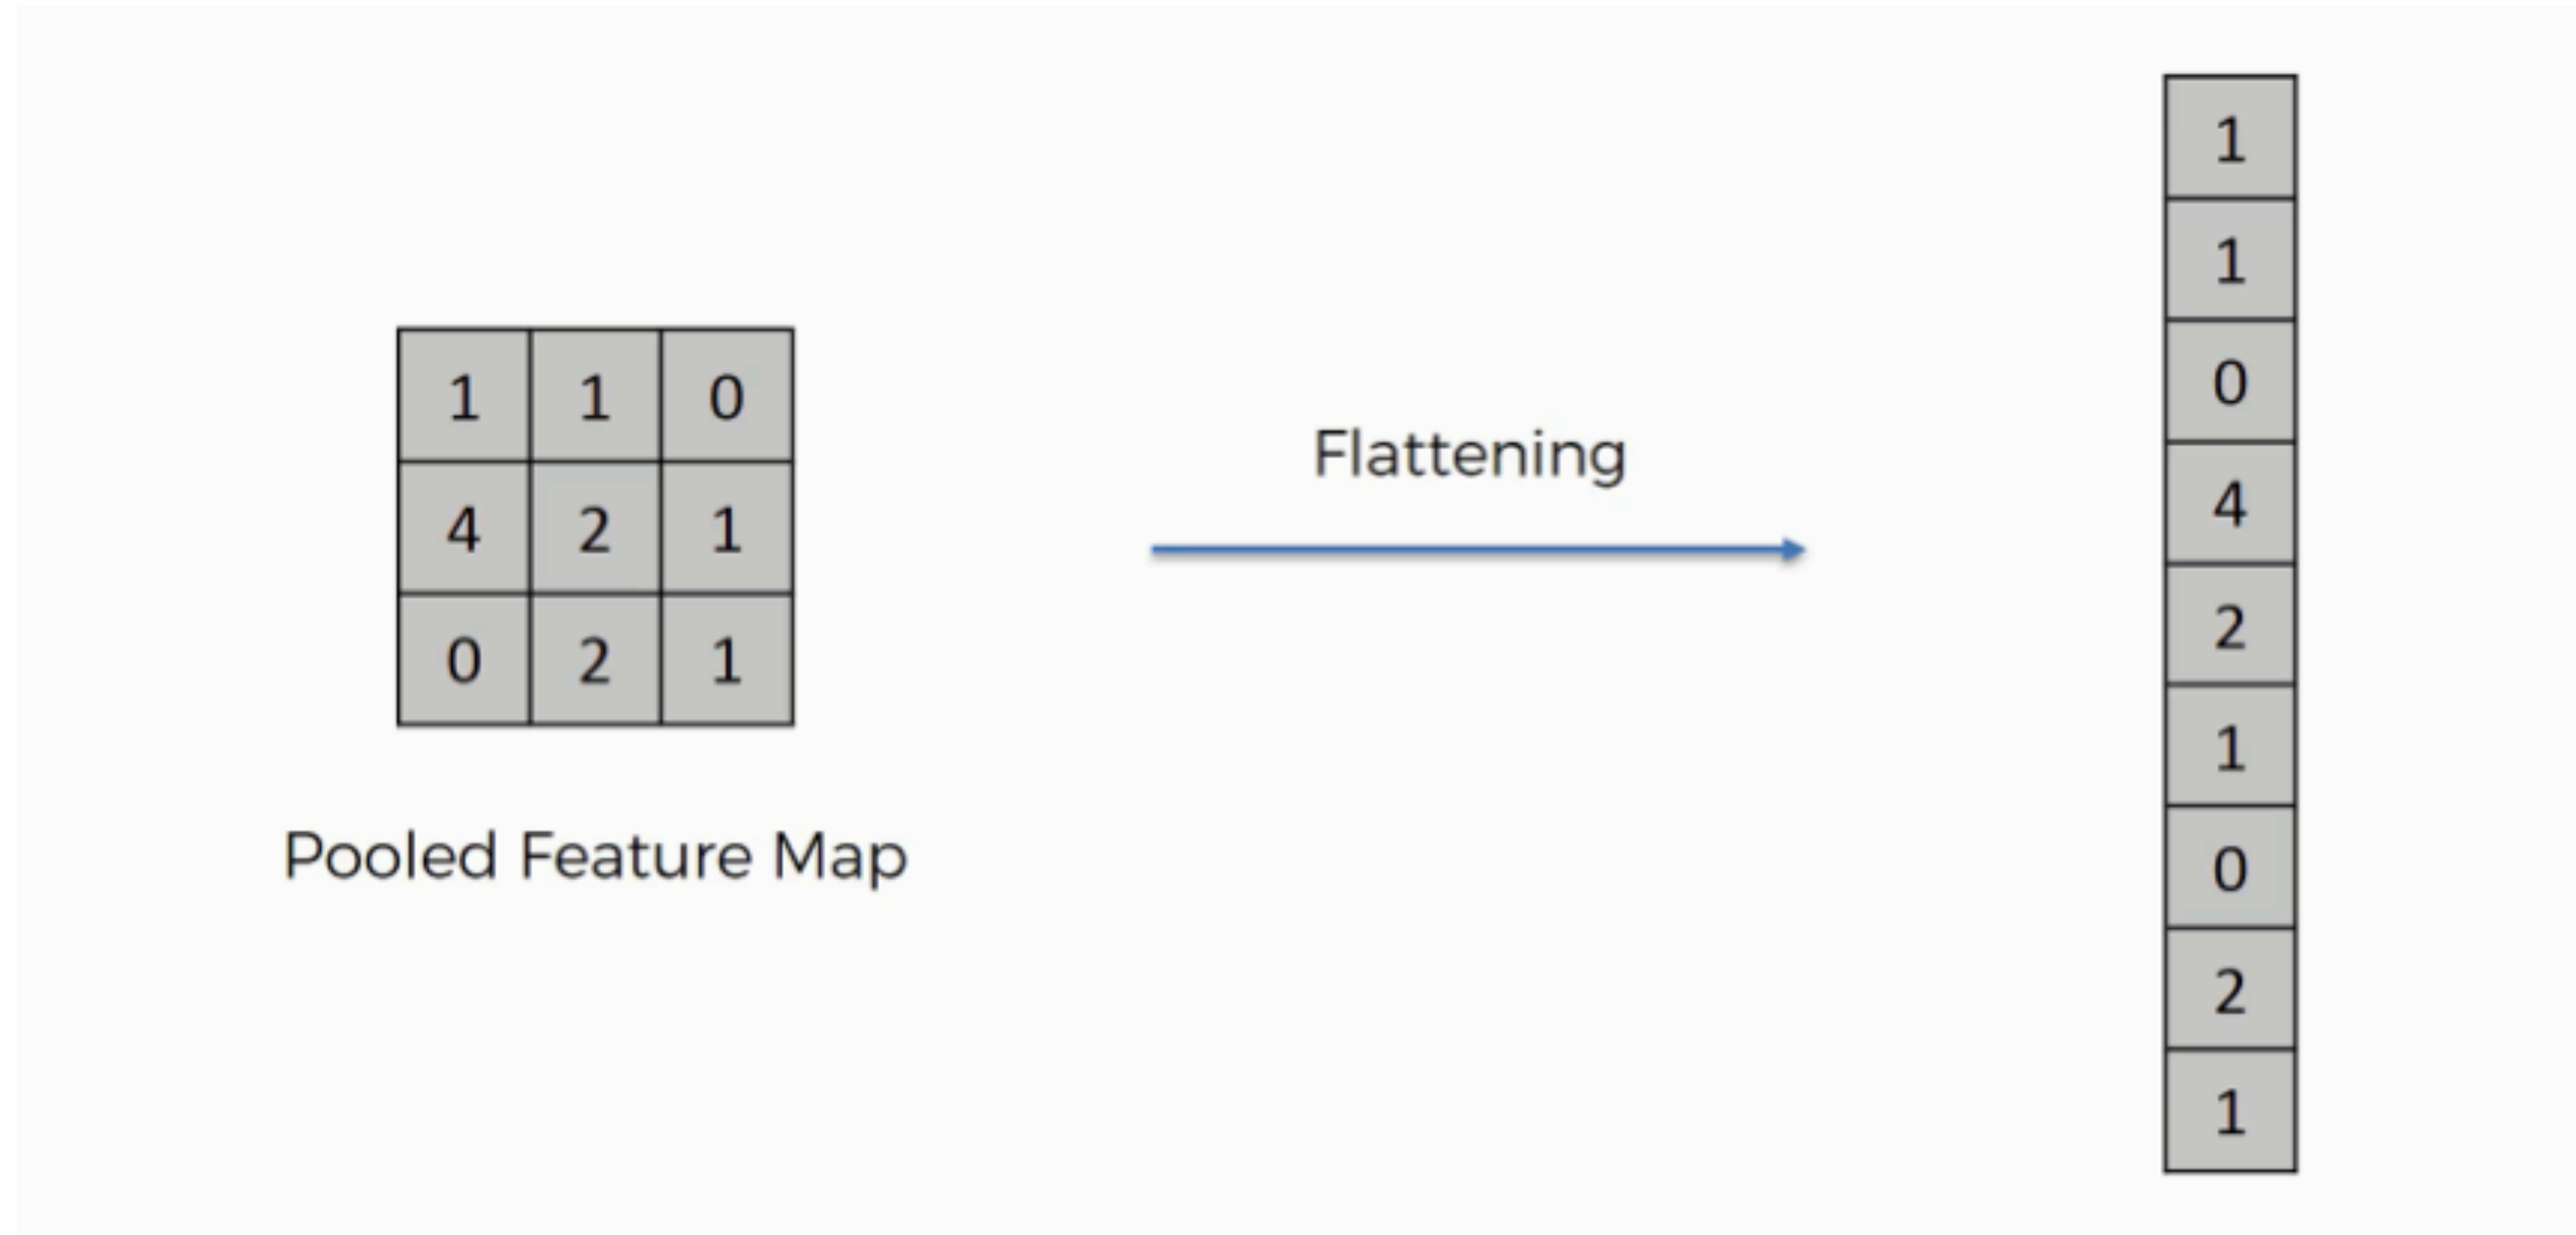
Return:

0	0	0	0
0	9	20	0

Flattening Layer

Purpose: Properly formats the pooled feature maps to be inserted in the dense layer

Takes dimension (X, Y, Z) pooled feature maps, and flattens them to an ($1, (X*Y*Z)$) flattened layer



Input Size = (1, 3, 3)

Flattened = (1 , (1*3*3)) => (1, 9)

Forward Propagation

```
self.before_flattened_shape = input.shape  
output = np.array([input.flatten()])
```

Back Propagation

```
before_flattened = input.reshape(self.before_flattened_shape)  
return before_flattened
```

Returns the flattened array back to its original shape

E.g. (1, 18) => (2, 3, 3)

Activation & Loss Functions

Binary Classification CNN:

“Classifying inputs into **two** categories”

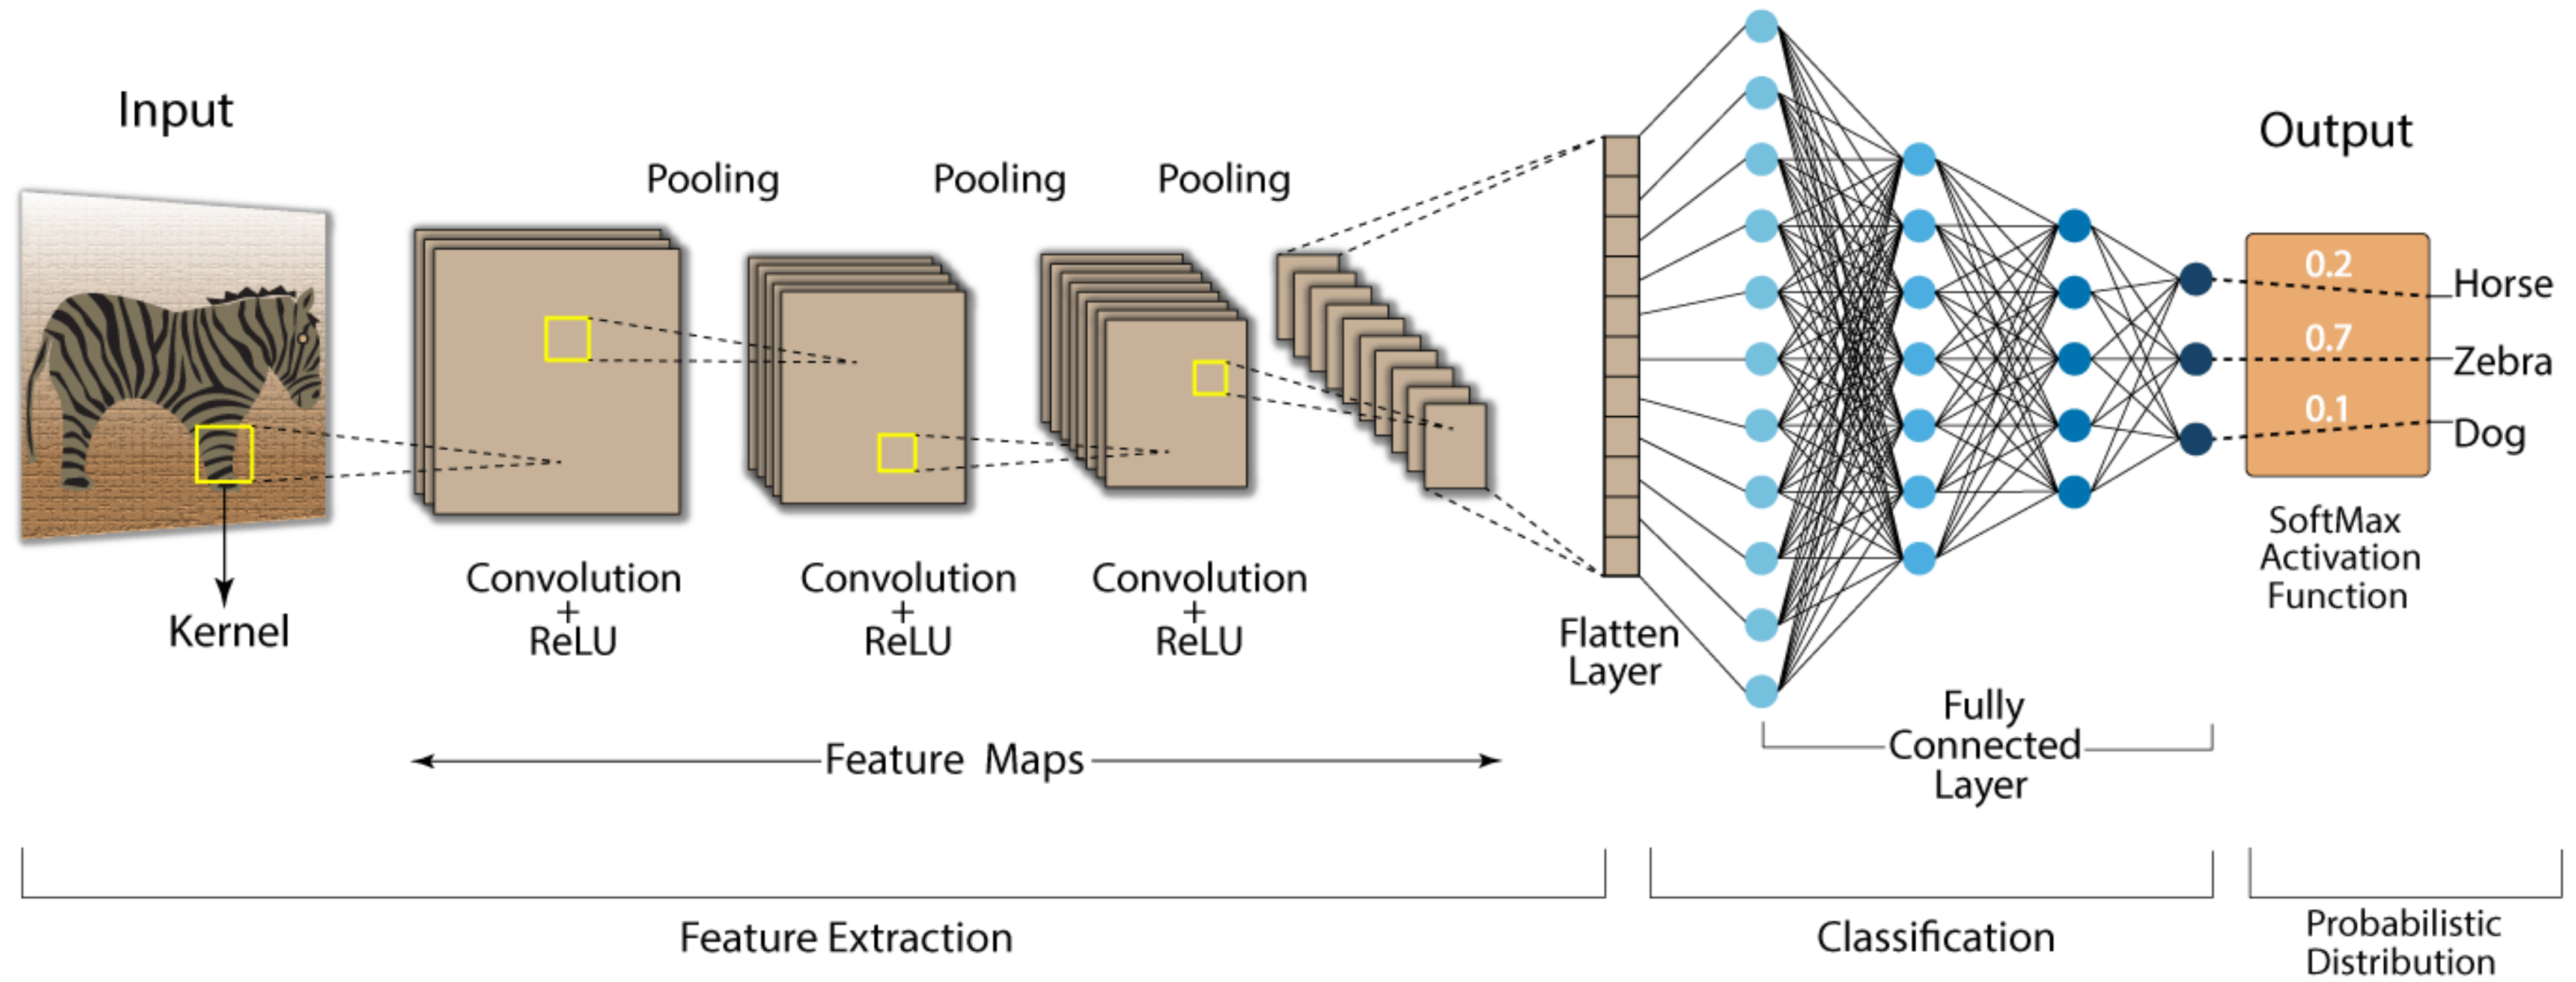
1. Activation = Sigmoid
2. Loss = Log Loss

Multi-Class Classification CNN:

“Classifying inputs into **multiple** categories”

1. Activation = Softmax
2. Loss = Categorical Cross Entropy

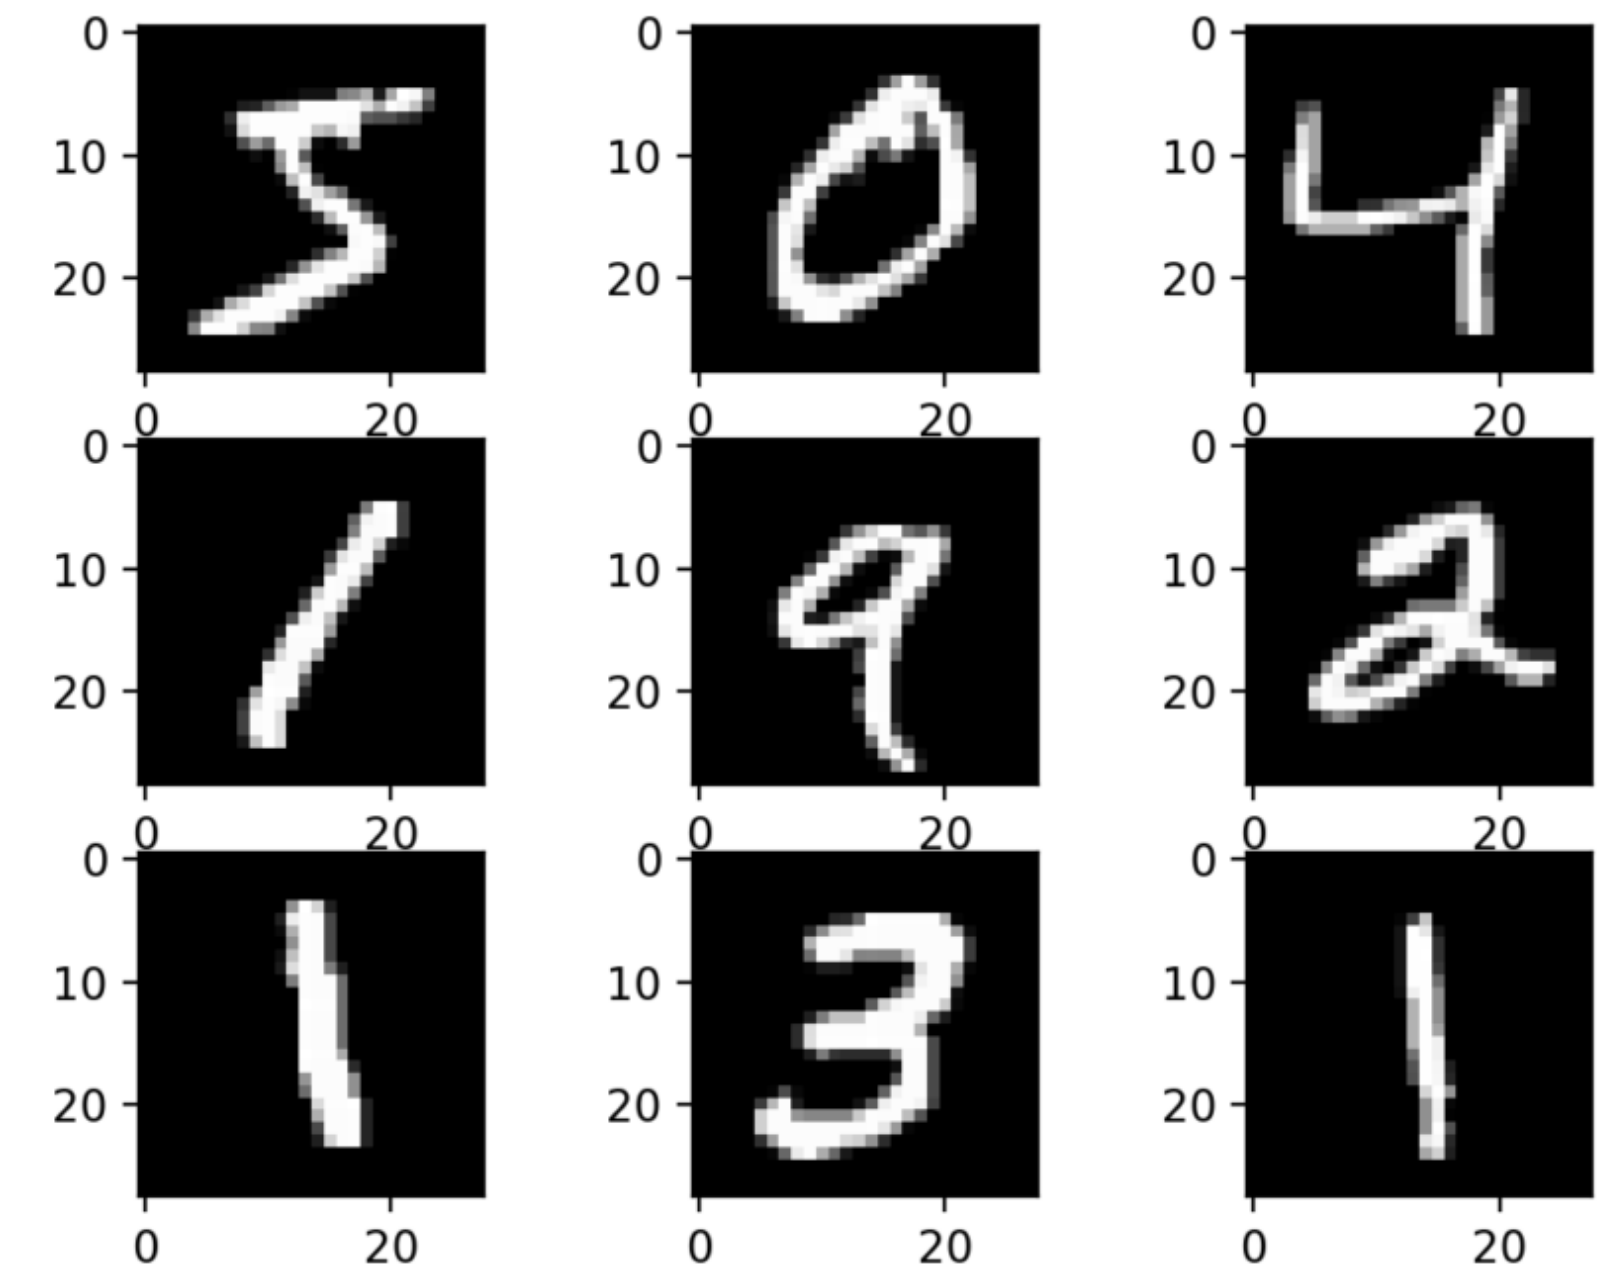
Convolution Neural Network (CNN)



Example

Data: **MNIST Handwritten Digit Dataset**

Implemented a **Binary** Classification CNN that identifies images (28 x 28) of 0 and 1 from the dataset



CNN Model (Binary Classification)

```
network = GWUNetwork()  
network.add(Conv2D(input_size=28, kernel_size=3))  
network.add(MaxPooling2D(pool_size=2, strides=2, input_size=23))  
network.add(Flatten(input_size=(11,11)))  
network.add(Dense(100, add_bias=False, activation='relu'))  
network.add(Dense(2, add_bias=False, activation='sigmoid'))  
network.compile(loss='log_loss', lr=0.001)  
network.fit(x_train_subset, y_train_subset, epochs=2)  
results = network.predict(x_test_subset)
```

```
# Import MNIST data
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train_subset = []
y_train_subset = []
x_test_subset = []
y_test_subset = []
```

```
# Separate images into two classes
for i in range(0, 1500):
    if (y_train[i] == 0 or y_train[i] == 1):
        x_train_subset.append(x_train[i])
        y_train_subset.append(y_train[i])
```

```
for i in range(0, 150):
    if (y_test[i] == 0 or y_test[i] == 1):
        x_test_subset.append(x_test[i])
        y_test_subset.append(y_test[i])
```

```
x_train_subset = np.array(x_train_subset[:200])
x_test_subset = np.array(x_test_subset[:10])
y_train_subset = np.array(y_train_subset[:200])
# Use keras to_categorical for binary classification problem
y_train_subset = np.array(tf.keras.utils.to_categorical(y_train_subset, num_classes=2))
y_test_subset = np.array(y_test_subset[:10])

# Normalize data
x_train_subset = x_train_subset.reshape(x_train_subset.shape[0], 28,
28).astype('float32')
x_test_subset = x_test_subset.reshape(x_test_subset.shape[0], 28, 28).astype('float32')
x_train_subset /= 255.0
x_test_subset /= 255.0
```

Output: Size 2 array that outputs the probability of the image being a “0” or a “1”

Example:

y_true [0, 1]




Output [0.2583, 0. 7417]

Output

```
[array([[0.5, 0.5]]), array([[0.52618988, 0.52502099]]), array([[0.5, 0.5]]), array([[0.5, 0.5]]),  
array([[0.48784266, 0.54074427]]), array([[0.5, 0.5]]), array([[0.57479947, 0.53996119]]), array(  
[[0.51681995, 0.51206615]]), array([[0.5059675 , 0.55003086]]), array([[0.5, 0.5]])]
```

```
Actual: [1, 0, 1, 0, 0, 1, 0, 0, 1, 1]
```

```
Predic: [1, 0, 1, 1, 1, 1, 0, 0, 1, 1]
```



Final Thoughts:

1. Learned a lot about the details of how a CNN works
2. Most difficult part personally was the back propagation of the convolutional layer and understanding the math behind it. Current implementation needs tweaking

References:

<https://www.youtube.com/watch?v=Lakz2MoHy6o>

<https://pavisj.medium.com/convolutions-and-backpropagations-46026a8f5d2c>

<https://towardsdatascience.com/backpropagation-in-a-convolutional-layer-24c8d64d8509>

<https://www.jefkine.com/general/2016/09/05/backpropagation-in-convolutional-neural-networks/>

Thank You.